

Reviewer Revision v2

Reproducible architecture · robust statistics · neutral prompts
atomic LLM execution · stability · grounding · surrogate fidelity

DESIGN DOCUMENT · NO RESULTS IMPLIED

Reviewer concern matrix

Concern	Status	Action
Train-only boundary	Technically closed	Hashed train-only summaries, demos, clustering and selection.
Heavy-tailed statistics	Technically closed	Untrimmed observations with P5/Q1/median/Q3/P95.
Stereotype priming	Technically closed	Neutral prompts and evidence/interpretation separation.
Cluster stability	Awaiting results	Seeds, thresholds, K, coverage, encoding, model matching.
Claim grounding	Awaiting results	Blinded judges, kappa, client bootstrap and adjudication.
Classifier faithfulness	Awaiting results	Surrogate fidelity plus classifier occlusion.
Human usefulness	Explicit limitation	No utility claim without a user study.

02 Found defects and corrections

Amount meanings, negative sorting, Rosbank recency, hidden reasoning leakage and label-filtered clusters are corrected and regression-tested.

03 Revision architecture

Versioned profiles feed neutral prompts, atomic generation, stable claims, label-agnostic clustering, frozen transforms and seeded evaluation.



04 Train-only data boundary

Validation selects prompts and hyperparameters. Validation/test labels never construct summaries, demonstrations, clusters, filters or features.



05 Artifact invalidation DAG

Prompt changes invalidate explanations, claims, embeddings, clusters, features and CoT/concat ML; independent baselines remain reusable.



06 Robust transaction statistics

Each dataset declares amount semantics; untrimmed client-level quantiles and category prevalence/share export to JSON, CSV, Markdown and LaTeX.

07 Neutral gender pilot

Four variants use 400 stratified validation clients (seed 137), balanced accuracy and a predeclared zero-shot equivalence rule.

08 Atomic API scheduler

Any 429 rolls back the whole window; 64 → 60-second cooldown → ten clean windows of 10 → probe 64.

HIGH 64

ROLLBACK + 60s

10 × 10

09 64 → cooldown → 10×10 → 64 state machine

Fallback counts only committed clean windows. A 429 at 10 rolls back the window and resets recovery.

HIGH 64

ROLLBACK + 60s

10 × 10

10 Parallel Qwen / GPT queues

Workers have independent semaphores, limiter states, locks, events and logs. The launcher is dry-run by default.

11 Live status and provenance

Append-only events drive terminal and HTML progress, error, concurrency, retry, throughput and checkpoint views.

12 Claims and semantic clustering

Stable claim IDs preserve source and extractor hashes; duplicate texts embed once; labels are post-hoc metadata only.

13 Publication cluster map

UMAP is presentation-only: atomic colors, prevalence-scaled points, gray outliers, macro-theme contours and callouts.

14 Stability experiments

Seeds 17/101/947; common 300-client anchors; thresholds, fixed K, coverage, encodings and ML seeds are reported as separate axes.

15 Grounding

Fifty clients and up to three claims per dataset/model cell; exact hashed evidence; two blinded judges; Codex adjudicates selected cases.

16 Classifier fidelity

A validation-selected non-claim teacher is compared with claim-based logistic, XGBoost and shallow-tree surrogates.

17 Cluster and tree explorer

Atomic clusters, macro-theme contours, representative claims, prevalence, compactness, lift, cross-model matches and tree paths.

18 Rerun matrix

Gender gets neutral API reruns; age and Rosbank retain immutable legacy API outputs and receive offline analyses; scoring is excluded.

19 Future launch commands

Computational runners require `--execute`; API runners additionally require `--execute-api` and `--until-complete`.

20 Reviewer response checklist

Stability, grounding and fidelity remain awaiting results; human usefulness remains an explicit limitation.